

Data Acquisition & Import

Sharing our Data for Replicable Science

The primary data source in this project was the **Million Playlist Dataset** provided by Spotify as part of the RecSys Challenge 2018. <https://labs.spotify.com/2018/05/30/introducing-the-million-playlist-dataset-and-recsys-challenge-2018/> (<https://labs.spotify.com/2018/05/30/introducing-the-million-playlist-dataset-and-recsys-challenge-2018/>) The challenge ended on June 30th of 2018, so we used an archival copy of the dataset provided to us by the course staff. All source code for this project is on a public GitHub Repository <https://github.com/IACS-CS-209-Group44/Spotify> (<https://github.com/IACS-CS-209-Group44/Spotify>)

Unfortunately, Git and GitHub are not well suited to working on large files, especially large binary files. Because of the large size of this data set, we were forced to use a shared folder on Dropbox to share large files. This makes it more challenging for us to share our data with the public and create a conveniently reproducible set of calculations. We are committed to the goal of fully reproducible science including data science with large files. For a full scale research undertaking with a suitable budget, two ideas to better achieve this goal would be to host a publicly available database instance using a service like AWS and / or to share a container instance. These techniques are beyond the scope of this project.

The data on our Dropbox includes several formats that make it easy to replicate almost all of our work. The h5 data files allow all of our imported data and results of calculations into Pandas dataframes. There is also a database backup available in the file `SpotifyDB.bak`. Warning: this is a bit of a monster, weighing in at **165 GB**, so it's not for the faint of heart. Someone with access to a moderately powerful SQL Server instance, either on-premises or on the cloud, can restore this database from the backup file. Anyone who is interested in accessing our dataset may email me at mse99@g.harvard.edu (<mailto:mse99@g.harvard.edu>) or michael.s.emanuel@gmail.com (<mailto:michael.s.emanuel@gmail.com>) to request access to this dataset. Any reasonable request will be granted read only access to the shared folder on Dropbox for as long as it is up.

The JSON Data Set and the Choice of SQL Server Backend

The data provided by the mpd consists of 1000 json files. Each file is a "slice" of 1,000 playlists. There are 1,000 of these slice files that together comprise the 1,000,000 playlists in the data set. The JSON files are a highly denormalized representation of the data, meaning they have a large amount of duplicated or redundant data. As one prominent example, each spotify track has a unique identifier, and a track name associated with this identifier. The JSON files duplicate the full track name when describing each track entry. We stored the data in a fully normalized format, with separate database tables for the logical entities of a Track, Playlist, PlaylistEntry, etc. I will describe our data model in greater detail below.

The choice to use SQL Server as our back end for data was an important strategic decision on this project. The JSON representation of the data is convenient for a person reviewing the data, but is extremely inefficient for doing large scale computations. As one example of how slow the JSON / Python API is, the data set comes with some Python utilities that do very simple summary calculations such as tabulating the most popular tracks. This utility was on pace to take over 90 minutes to run on my desktop PC, which has top of the line hardware for a desktop computer.

It was clear that we needed a far more efficient back end for storing our data in a more normalized form. Most people in this course would probably have chosen to use Pandas and persisted a series of data frames. That is a completely sound choice. In my particular case, I have spent many years working with SQL (first MySQL, then SQL Server). I have also invested a large amount of time and money configuring an instance of SQL Server running on a high performance server sitting on a rack in my basement. By comparison, I am very new to Pandas. I'm now proficient for basic tasks, but I have often spent hours trying to figure out how to do an operation I could do in a matter of minutes in SQL Server. It is often true that the best tool for a given job is the one you know how to use. For me and this problem, that tool was SQL Server.

Data Import into SQL Server

The first step was to understand the data and break it up into logical entities. Each logical entity corresponds to its own database table. The four most important logical entities in this data set are Artist, Album, Track, and Playlist. I am going to go out of order for a moment and jump ahead to the Pandas dataframes that we used for downstream computations. The code below loads these frames from an h5 data file.

Before this code will run, please copy the files `data.h5` and `playlist_entry.h5` from the Dropbox folder to the directory where you cloned the GitHub repo. On my system, the repo is cloned to `D:\IACS-CS-209-Spotify\` and the Dropbox folder is at `D:\Dropbox\IACS-CS-209-Spotify\`. I copy the file `D:\Dropbox\IACS-CS-209-Spotify\mpd\database_export\h5` into

D:\IACS-CS-209-Spotify\data\data.h5 . I do the analogous operation for files `playlist_entry.h5` and `track_pair.h5` .

It is possible to automate all of this in a very slick and convenient way using GitHub LFS. <https://git-lfs.github.com/> (<https://git-lfs.github.com/>) But this is a paid service that costs at least \$10 a month, possibly quite a bit more given the size of this data set. I spent several hours investigating alternatives that would allow us to bypass this manual copy step to synchronize the Dropbox folder, but did not come up with a better procedure. Once the file `data.h5` is in place, the code below will load the data frames for Artist, Album, Track, and Playlist into memory.

```
In [1]: import pandas as pd
import time
from IPython.display import display
from typing import List, Dict, Optional
```

```

In [2]: def load_frames(frame_names: Optional[List[str]] = None) -> Dict[str, pd.DataFrame]:
        """Load all available data frames. Return a dictionary keyed by frame_name."""
        # Relative path to h5 data files
        path_h5 = '../data/'

        # Dictionary of dataframes to be generated.
        # Key = frame name, value = fname_h5
        frame_tbl: Dict[str, str] = {
            # Basic schema
            'Artist': 'data.h5',
            'Album': 'data.h5',
            'Track': 'data.h5',
            'Playlist': 'data.h5',

            # Tables relating to prediction outcomes and scoring
            'TrainTestSplit': 'data.h5',
            'Playlist_Last10': 'data.h5',
            'Playlist_trn': 'data.h5',
            'Playlist_tst': 'data.h5',

            # Tables relating to the baseline and playlist name prediction models
            'TrackRank': 'data.h5',
            'PlaylistName': 'data.h5',
            'PlaylistSimpleName': 'data.h5',
            'TrackRankBySimpleName': 'data.h5',

            # PlaylistEntry table is big - saved in its own file
            'PlaylistEntry': 'playlist_entry.h5',

            # Audio features
            'AudioFeatures': 'data.h5',
            'Genre': 'data.h5',
            'MetaGenre': 'data.h5',
            'TrackGenre': 'data.h5',
            'TrackMetaGenre': 'data.h5',

            # TrackPairs table is big - saved in its own file
            'TrackPairs': 'track_pairs.h5',

            # Scores of three models: baseline, playlist name, naive bayes
            'Scores_Baseline': 'data.h5',
            'Scores_SimpleName': 'data.h5',
            'Scores_TrackPair': 'data.h5',
            'Scores_Stack': 'data.h5',

            # Survey responses
            'SurveyResponse': 'data.h5',
            'SurveyPlaylist': 'data.h5',
            'SurveyPlaylistEntry': 'data.h5',

            # Artists being promoted by policy (mid-tier, female)
            'PromotedArtist': 'data.h5',

            # Survey recommendations
            'SurveyRecommendations': 'data.h5',
            'SurveyRecommendationsPromoted': 'data.h5',
        }

        # Set frame_names to all tables if it was not specified
        if frame_names is None:
            frame_names = frame_tbl.keys()

        # Start timer
        t0 = time.time()
        # Dictionary of data frames
        frames: Dict[str, pd.DataFrame] = dict()
        # Iterate over entries in frame_names, loading them from h5 files
        for frame_name in frame_names:
            # h5 filename for this frame
            fname_h5 = frame_tbl[frame_name]
            # Read the data frame

```

```

frames[frame_name] = pd.read_hdf(path_h5 + fname_h5, frame_name)
# Status update
print(f'Loaded {frame_name}.')

# Status update
t1 = time.time()
elapsed = t1 - t0
print(f'\nLoaded {len(frames)} Data Frames.')
print(f'Elapsed Time: {elapsed:0.2f} seconds.')
return frames

```

Load the frames for Artist, Album, Track, Playlist & PlaylistEntry into memory

```
In [9]: frames = load_frames(['Artist', 'Album', 'Track', 'Playlist', 'PlaylistEntry'])
```

```

Loaded Artist.
Loaded Album.
Loaded Track.
Loaded Playlist.
Loaded PlaylistEntry.

Loaded 5 Data Frames.
Elapsed Time: 4.72 seconds.

```

Notice how much faster and more efficient this is than loading the JSON files. An analogous operation run directly on the JSON files took on the order of multiple minutes on my system. The size of the decompressed JSON files was approximately 32 GB. The size of the normalized tables storing the exact same information was approximately 2 GB. Normalizing the data compressed it by a factor of 16. (Try doing that with a ZIP file compression!)

It is said that a picture is worth a thousand words. The same is often true of an example. Please see below examples of the first 10 records on the five key tables in the data set: Artist, Album, Track, Playlist, and PlaylistEntry.

Data Model for Artist, Album, Track, Playlist & PlaylistEntry

```
In [4]: display(frames['Artist'].head(10))
```

	ArtistID	ArtistUri	ArtistName
0	1	spotify:artist:0001cekkfdEBoMlwVQvpLg	Jordan Colle
1	2	spotify:artist:0001wHqxbF2YYRQxGdbyER	Motion Drive
2	3	spotify:artist:0001ZVMPt41Vwzt1zsmuzp	Thyro & Yumi
3	4	spotify:artist:0004C5XZIKZyd2RWvP4sOq	"Faron Young, Nat Stuckey"
4	5	spotify:artist:000DnGPNOsxvqb2YEHBePR	The Ruins
5	6	spotify:artist:000Dq0VqTZpxOP6jQMscVL	Thug Brothers
6	7	spotify:artist:000h2XLY65iWC9u5zgcL1M	Kosmose
7	8	spotify:artist:000spuc3oKgwYmfg5IE26s	Parliament Syndicate
8	9	spotify:artist:000UUAIdQqkTD9sfoyQGf	Darren Gibson
9	10	spotify:artist:000UxvYLQuybj6iVRRCAw1	Primera Etica

```
In [5]: display(frames['Album'].head(10))
```

	AlbumID	AlbumUri	AlbumName
0	1	spotify:album:00010fh2pSk7f1mGlhgorB	Okkadu (Original Motion Picture Soundtrack)
1	2	spotify:album:00045VFusrXwCSietfmspc	Let Love Begin Remixed
2	3	spotify:album:0005lpYtyKk9B3e0mWJdem	Stability
3	4	spotify:album:0005rH90S3le891y5XzPg4	Mozart: Piano Concerto No. 27, KV595
4	5	spotify:album:0008WZMLnvEBVnq418uZsl	Smart Flesh
5	6	spotify:album:0009lq7uJ6cW3Cxtf8eNUp	Earth: The Pale Blue Dot (Instrumental)
6	7	spotify:album:000aG92zPFtZ0FRLaaJHE5	X
7	8	spotify:album:000f3dTvpazVzv35NuZmn	Make It Fast, Make It Slow
8	9	spotify:album:000g9ysmwb8NNsd4u1o087	Nennt es, wie Ihr wollt
9	10	spotify:album:000gdWY9uR4VYS5oZudY5o	Pérez Prado. Sus 40 Grandes Canciones

```
In [16]: display(frames['Track'].head(10))
```

	TrackID	ArtistID	AlbumID	TrackUri	TrackName
0	1	208716	355266	spotify:track:0000uJA4xCdxThagdLkkLR	Heart As Cold As Stone
1	2	110598	6666	spotify:track:0002yNGLtYSYtc0X6ZnFvp	Muskrat Ramble
2	3	93681	77525	spotify:track:00039MgrmLolzSpuYKurn9	Thas What I Do
3	4	5377	586855	spotify:track:0003Z98F6hUq7XxqSRM87H	???? ?????? ??? ???
4	5	285766	426742	spotify:track:0004ExljAge0P5XWn1LXmW	Gita
5	6	240510	337666	spotify:track:0005rgjsSeVLp1cze57jIN	Mi Razón de Ser
6	7	111429	625930	spotify:track:0005w1bMJ7QAMl6DY98oxa	Sonata in G Major, BuxWV 271: Allegro -
7	8	102394	113691	spotify:track:0006Rv1e2Xfh6QooyKJqKS	Nightwood
8	9	240180	287678	spotify:track:0007AYhg2UQbEm88mxu7js	Mandarin Oranges Part 2
9	10	35528	570255	spotify:track:0009mEWM7HILVo4VZYtqwc	Movement

```
In [17]: display(frames['Playlist'].head(10))
```

	PlaylistID	PlaylistName	NumTracks	NumArtists	NumFollowers	NumEdits	DurationMS	IsCollaborative	ModifiedAt
0	0	Throwbacks	52	37	1	6	11532414	0	1493424000
1	1	Awesome Playlist	39	21	1	5	11656470	0	1506556800
2	2	korean	64	31	1	18	14039958	0	1505692800
3	3	mat	126	86	1	4	28926058	0	1501027200
4	4	90s	17	16	2	7	4335282	0	1401667200
5	5	Wedding	80	56	1	3	19156557	0	1430956800
6	6	I Put A Spell On You	16	13	1	2	3408479	0	1477094400
7	7	2017	53	48	1	38	12674796	0	1509321600
8	8	BOP	46	23	2	21	9948921	0	1508976000
9	9	old country	21	18	1	10	4297488	0	1501804800

```
In [18]: display(frames['PlaylistEntry'].head(10))
```

	PlaylistID	Position	TrackID
0	0	0	236619
1	0	1	1866537
2	0	2	260403
3	0	3	347127
4	0	4	451364
5	0	5	270971
6	0	6	1784688
7	0	7	938244
8	0	8	2145897
9	0	9	776743

Comments on the Table Design

The design of these database tables follows a few simple best practices. All tables have an integer ID as their primary key. This gives a large boost to the performance of queries that join tables. Searching for an integer entry in an index is a much faster operation than comparing strings. The original JSON data model did not have any integer IDs for any of these entities besides for the playlist ID. All entities that exist in Spotify also have a field that Spotify names a Uri. These are string identifiers. These fields are equipped with unique constraints, building in both a data integrity check and causing SQL server to build indexes that both enforce the constraint and support fast joins on these fields.

The Track table demonstrates foreign key relationships. The fields ArtistID and AlbumID are foreign keys onto the Artist and Album tables, respectively. Note that no redundant information such as the ArtistName or ArtistUri are stored in the Track table. Any consumer of this information is expected to get it by joining the Artist table using ArtistID. The foreign key relationships are enforced with foreign key constraints. The primary key on the Track table is TrackID, and there is a separate unique constraint on the TrackUri.

This table also demonstrates a consistent naming scheme followed in the SpotifyDB database. The name of the integer primary key on the table Artist is ArtistID. This is one popular approach. Another approach is to name the field ID. I prefer to name the field ArtistID because then when you join from the Track table to the Artist table, the join clause uses the field ArtistID on both sides of the equality. There are multiple naming approaches that are strong. As is often the case, the most important thing is to pick one strategy and then follow it consistently.

For all you SQL aficionados out there, below please find the SQL table definitions for the main logical entities comprising the MPD dataset. (Don't worry, I won't include all 42 tables in the database, just a few of the important ones!). The entirety of the SQL used in this project can be found in the `sql` directory under the GitHub repo. The files are named with a numeric prefix, so that tables are built in the correct order. As an example, the Track table references the Artist and Album tables, so Artist and Album must be created first. The relevant SQL scripts are named `03_MakeTable_Artist.sql`, `04_MakeTable_Album.sql`, and `05_MakeTable_Track.sql`.

SQL Table Definitions for Main Logical Entities

Artist / Script `03_MakeTable_Artist.sql`

```

DROP sequence IF EXISTS dbo.SEQ_ArtistID
CREATE sequence dbo.SEQ_ArtistID
    AS INT start WITH 1 increment BY 1 NO cycle;

DROP TABLE IF EXISTS dbo.Artist;
CREATE TABLE dbo.Artist(
    ArtistID INT NOT NULL
        DEFAULT next value FOR dbo.SEQ_ArtistID,
    ArtistUri CHAR(37) NOT NULL,
    ArtistName VARCHAR(512) NOT NULL,
    -- Primary Key and Unique constraints
    CONSTRAINT PK_Artist_ArtistID PRIMARY KEY (ArtistID),
    CONSTRAINT UNQ_Artist_ArtistUri UNIQUE(ArtistUri),
    -- ArtistName should be unique, but unfortunately it's not; index it instead
    INDEX IDX_Artist_ArtistName (ArtistName),
);

```

Album / Script 04_MakeTable_Album.sql

```

DROP sequence IF EXISTS dbo.SEQ_AlbumID
CREATE sequence dbo.SEQ_AlbumID
    AS INT start WITH 1 increment BY 1 NO cycle;

DROP TABLE IF EXISTS dbo.Album;
CREATE TABLE dbo.Album(
    AlbumID INT NOT NULL
        DEFAULT next value FOR dbo.SEQ_AlbumID,
    AlbumUri CHAR(36) NOT NULL,
    AlbumName VARCHAR(512) NOT NULL,
    -- Primary Key and Unique constraints
    CONSTRAINT PK_Album_AlbumID PRIMARY KEY (AlbumID),
    CONSTRAINT UNQ_Album_AlbumUri UNIQUE(AlbumUri),
    -- AlbumName should be unique, but unfortunately it's not; index it instead
    INDEX IDX_Album_AlbumName (AlbumName),
);

```

Track / Script 05_MakeTable_Track.sql

```

DROP sequence IF EXISTS dbo.SEQ_TrackID
CREATE sequence dbo.SEQ_TrackID
    AS INT start WITH 1 increment BY 1 NO cycle;

DROP TABLE IF EXISTS dbo.Track;
CREATE TABLE dbo.Track(
    TrackID INT NOT NULL
        DEFAULT next value FOR dbo.SEQ_TrackID,
    ArtistID INT NOT NULL,
    AlbumID INT NOT NULL,
    TrackUri CHAR(36) NOT NULL,
    TrackName VARCHAR(512) NOT NULL,
    -- Primary Key and Unique constraints
    CONSTRAINT PK_Track_TrackID PRIMARY KEY (TrackID),
    CONSTRAINT UNQ_Track_TrackUri UNIQUE(TrackUri),
    -- TrackName should be unique, but unfortunately it's not; index it instead
    INDEX IDX_Track_TrackName (TrackName),
    -- Foreign keys on ArtistID and AlbumID
    CONSTRAINT FK_Track_ArtistID FOREIGN KEY (ArtistID)
        REFERENCES dbo.Artist(ArtistID),
    CONSTRAINT FK_Track_AlbumID FOREIGN KEY (AlbumID)
        REFERENCES dbo.Album(AlbumID)
);

```

Playlist / Script 06_MakeTable_Playlist.sql

```

CREATE TABLE dbo.Playlist(
    PlaylistID INT NOT NULL,
    PlaylistName varchar(100) NOT NULL,
    -- Playlist attributes
    NumTracks SMALLINT NOT NULL,
    NumArtists SMALLINT NOT NULL,
    NumFollowers INT NOT NULL,
    NumEdits SMALLINT NOT NULL,
    DurationMS INT NOT NULL,
    IsCollaborative BIT NOT NULL,
    ModifiedAt INT NOT NULL,
    -- Primary key
    CONSTRAINT PK_Playlist_PlaylistID PRIMARY KEY (PlaylistID),
    -- Indices
    INDEX IDX_Playlist_PlaylistName (PlaylistName),
    INDEX IDX_Playlist_NumFollowers(NumFollowers),
);

```

PlaylistEntry / Script 07_MakeTable_PlaylistEntry.sql

```

CREATE TABLE dbo.PlaylistEntry(
    PlaylistID INT NOT NULL,
    Position SMALLINT NOT NULL,
    TrackID INT NOT NULL,
    -- Primary key is the pair of the PlaylistID and Position
    CONSTRAINT PK_PlaylistEntry PRIMARY KEY (PlaylistID, Position),
    -- Foreign keys
    CONSTRAINT FK_PlaylistEntry_PlaylistID
        FOREIGN KEY (PlaylistID) REFERENCES dbo.Playlist(PlaylistID),
    CONSTRAINT FK_PlaylistEntry_TrackID
        FOREIGN KEY (TrackID) REFERENCES dbo.Track(TrackID),
);

```


Importing the Raw Data into SQL Server

These empty tables are very nice and clean, but they don't look anything like the quite bloated JSON files. How were they populated? I used a two step process. In the first step, I populated "raw" tables for Playlist and PlaylistEntry that followed the structure of the JSON file contents much more closely. These are the definitions of the raw tables; they have the same table names as their sister tables in the main database, but are stored in a different schema `r` (for raw) rather than the gnominically named schema `dbo` which is the default schema in SQL Server, and short for DataBaseOwner.

Raw Table `r.Playlist` / Script `01_MakeTable_r_Playlist`

```
CREATE TABLE r.Playlist(  
  PlaylistID INT NOT NULL,  
  PlaylistName VARCHAR(1024) NOT NULL,  
  -- The number of tracks, albums, and artists on this playlist  
  NumTracks SMALLINT NOT NULL,  
  NumAlbums SMALLINT NOT NULL,  
  NumArtists SMALLINT NOT NULL,  
  -- Additional information about the playlist  
  NumFollowers INT NOT NULL,  
  NumEdits SMALLINT NOT NULL,  
  DurationMS INT NOT NULL,  
  IsCollaborative BIT NOT NULL,  
  ModifiedAt INT NOT NULL,  
  -- Primary Key and indices  
  CONSTRAINT PK_r_Playlist_PlaylistID PRIMARY KEY (PlaylistID),  
  INDEX IDX_r_Playlist_PlaylistName (PlaylistName)  
)
```

Raw Table `r.PlaylistEntry` / Script `02_MakeTable_r_PlaylistEntry`

```
CREATE TABLE r.PlaylistEntry(  
  PlaylistID INT NOT NULL,  
  Position SMALLINT NOT NULL,  
  -- The track  
  TrackUri VARCHAR(256) NOT NULL,  
  TrackName VARCHAR(1024) NOT NULL,  
  -- The album  
  AlbumUri VARCHAR(256) NOT NULL,  
  AlbumName VARCHAR(1024) NOT NULL,  
  -- The artist  
  ArtistUri VARCHAR(256) NOT NULL,  
  ArtistName VARCHAR(1024) NOT NULL,  
  -- Additional track info  
  TrackDurationMS INT NOT null  
  -- Primary Key and indices  
  CONSTRAINT PK_r_PlaylistEntry PRIMARY KEY (PlaylistID, Position),  
  INDEX IDX_r_PlaylistEntry_Tracks (TrackUri, TrackName),  
  INDEX IDX_r_PlaylistEntry_Albums (AlbumUri, AlbumName),  
  INDEX idx_r_PlaylistEntry_Artists (ArtistUri, ArtistName)  
)
```

The cells below present the contents of the file `src/db_import.py`. This program was run from the terminal and it populated the two raw tables `r.Playlist` and `r.PlaylistEntry` by reading the JSON files. In practice, if you are replicating our results the best strategy is probably to restore database from `SpotifyDB.bak`. In principle, someone trying to replicate our work could run this script on their local system. They would need to modify the function `getConnection()` though because it is currently configured to acquire a connection to my database instance, which is on a server named `Thor` and an instance named `Mjolnir`. (Of course,

if you are a fan of Norse mythology or the Marvel Universe, you might *already* have a SQL Server Instance with this name, but it seems... unlikely.) Just change `Thor` to the name of your SQL Server and `Mjolnir` to the name of your instance and it should work with an on-premises setup. I don't know the exact steps to configure it on AWS but it should be straightforward.

```

In [20]: from sys import argv
import os
import json
import pyodbc
import time
from typing import List, Tuple, Dict

# *****
# Mapping from strings to bool for IsCollaborative field
str2bool: Dict[str, bool] = {
    'true': True,
    'false': False,
}

def get_insertions(fname: str) -> Tuple[List[Tuple], List[Tuple]]:
    """
    Read the mpd slice with this filename.
    Returns two lists of tuples, inserts_playlist and inserts_tracks.
    Each entry is a tuple matching one record on the r.Playlist and r.PlaylistEntry tables, respectively.
    """
    # Reference external variable used
    global str2bool
    # Open the file
    with open(fname) as fh:
        # Read the json contents
        js = fh.read()
    # Read in the data as a JSON object
    mpd_slice = json.loads(js)
    # Extract the playlists field from the slice
    playlists = mpd_slice['playlists']
    # Length (should be 1000)
    playlist_count: int = len(playlists)

    # Preallocate list of rows to be inserted for the Playlist table
    rows_playlist: List[Tuple] = playlist_count * [None]
    # Initialize an empty list of rows to be inserted for the PlaylistEntry table (we don't know its length)
    rows_playlist_entry: List[Tuple] = list()

    # Iterate over each playlist in the slice
    for i, playlist in enumerate(playlists):
        # Get the attributes of this playlist
        # Name attributes consistent with the database naming scheme
        # ID and name
        PlaylistID: int = playlist['pid']
        PlaylistName: str = playlist['name']
        # Number of tracks, albums, and artists
        NumTracks: int = playlist['num_tracks']
        NumAlbums: int = playlist['num_albums']
        NumArtists: int = playlist['num_artists']
        # Additional info
        NumFollowers: int = playlist['num_followers']
        NumEdits: int = playlist['num_edits']
        DurationMS: int = playlist['duration_ms']
        IsCollaborative: bool = str2bool[playlist['collaborative']]
        ModifiedAt: int = playlist['modified_at']

        # Assemble this into a tuple with one row to be inserted into r.Playlist
        row_playlist: Tuple = (PlaylistID, PlaylistName, NumTracks, NumAlbums, NumArtists,
                               NumFollowers, NumEdits, DurationMS, IsCollaborative, ModifiedAt)
        # Save this row to the inserts
        rows_playlist[i] = row_playlist

        # Get the tracks out of this playlist
        tracks: List[Dict] = playlist['tracks']
        # Iterate over the tracks
        for track in tracks:
            # Get the contents of this track - use database names and order
            # already have PlaylistID above
            Position: int = track['pos']

```

```

# The track
TrackUri: str = track['track_uri']
TrackName: str = track['track_name']
# The album
AlbumUri: str = track['album_uri']
AlbumName: str = track['album_name']
# The artist
ArtistUri: str = track['artist_uri']
ArtistName: str = track['artist_name']
# Duration
TrackDurationMS: int = track['duration_ms']

# Assemble this into a tuple with one row to be inserted into r.PlaylistEntry
row_playlist_entry: Tuple = (PlaylistID, Position, TrackUri, TrackName, AlbumUri, AlbumName,
                             ArtistUri, ArtistName, TrackDurationMS)
rows_playlist_entry.append(row_playlist_entry)

# Return the lists ready to be inserted into r.Playlist and r.PlaylistEntry
return (rows_playlist, rows_playlist_entry)

```

```

In [23]: # *****
def getConnection() -> pyodbc.Connection:
    """Get database connection"""
    raise RuntimeError('Oops please do not run this again the table is already populated!')
    driver: str = r'{ODBC Driver 13 for SQL Server}'
    server: str = r'THOR\MJOLNIR'
    database: str = 'SpotifyDB'
    auth: str = 'Trusted_Connection=yes;'
    conn_string: str = f'DRIVER={driver};SERVER={server};DATABASE={database};{auth}'
    conn: pyodbc.Connection = pyodbc.connect(conn_string)
    return conn

# *****
def delete_playlist(curs, PlaylistID_Min: int, PlaylistID_Max: int) -> None:
    """
    Deletes a block of rows in DB table r.Playlist
    INPUTS:
    =====
    curs:          Database cursor
    PlaylistID_min: First PlaylistID to be deleted (inclusive)
    PlaylistID_max: Last  PlaylistID to be deleted (exclusive)
    """

    # SQL string to delete records for this range of PlaylistID
    sqlDelete = '''
    DELETE FROM r.Playlist WHERE ? <= PlaylistID and PlaylistID < ?
    '''

    # Delete records in this range of PlaylistID
    curs.execute(sqlDelete, PlaylistID_Min, PlaylistID_Max)

def insert_playlist(curs, rows: List[Tuple]):
    """
    Inserts a list of rows into the DB table r.Playlist
    INPUTS:
    =====
    curs:          Database cursor
    rows_playlist: Row of records to be inserted
    """

    # SQL string to insert ONE record into r.Playlist
    # row_playlist: Tuple = (PlaylistID, PlaylistName, NumTracks, NumAlbums, NumArtists,
    #                        NumFollowers, NumEdits, DurationMS, IsCollaborative, ModifiedAt)
    sqlInsert = '''
    INSERT INTO r.Playlist
    (PlaylistID, PlaylistName, NumTracks, NumAlbums, NumArtists,
    NumFollowers, NumEdits, DurationMS, IsCollaborative, ModifiedAt)
    VALUES (?, ?, ?, ?, ?, ?, ?, ?, ?, ?)
    '''

    # Insert batch of records using executemany method
    curs.executemany(sqlInsert, rows)
    # Commit changes
    curs.commit()

# *****
def delete_playlist_entry(curs, PlaylistID_Min: int, PlaylistID_Max: int) -> None:
    """
    Deletes a block of rows in DB table r.PlaylistEntry
    INPUTS:
    =====
    curs:          Database cursor
    PlaylistID_min: First PlaylistID to be deleted (inclusive)
    PlaylistID_max: Last  PlaylistID to be deleted (exclusive)
    """

    # SQL string to delete records for this range of PlaylistID
    sqlDelete = '''
    DELETE FROM r.PlaylistEntry WHERE ? <= PlaylistID and PlaylistID < ?
    '''

```

```

# Delete records in this range of PlaylistID
curs.execute(sqlDelete, PlaylistID_Min, PlaylistID_Max)

def insert_playlist_entry(curs, rows: List[Tuple]):
    """
    Inserts a list of rows into the DB table r.PlaylistEntry
    INPUTS:
    =====
    curs:          Database cursor
    rows_playlist: Row of records to be inserted
    """

    # SQL string to insert ONE record into r.PlaylistEntry
    # row_playlist_entry: Tuple = (PlaylistID, Position, TrackUri, TrackName, AlbumUri, AlbumName,
    #                               ArtistUri, ArtistName, TrackDurationMS)
    sqlInsert = '''
    INSERT INTO r.PlaylistEntry
    (PlaylistID, Position, TrackUri, TrackName, AlbumUri, AlbumName, ArtistUri, ArtistName, TrackDurationM
    VALUES (?, ?, ?, ?, ?, ?, ?, ?, ?)
    '''

    # Insert batch of records using executemany method
    curs.executemany(sqlInsert, rows)
    # Commit changes
    curs.commit()

```

```

In [26]: # *****
def main():
    # Unpack arguments
    argc: int = len(argv)-1
    if argc == 0:
        sliceMin = 0
        sliceMax = 1000
    if argc == 1:
        sliceMin = 0
        sliceMax = int(argv[1])
    if argc == 2:
        sliceMin: int = int(argv[1])
        sliceMax: int = int(argv[2])
    if argc not in (0, 1, 2):
        print('Usage: python db_import.py sliceMin sliceMax.')
        print('This will insert slices from and including sliceMin, up to but not including sliceMax.')
        return
    # Range of PlaylistID's
    PlaylistID_Min: int = sliceMin * 1000
    PlaylistID_Max: int = sliceMax * 1000
    # Status update
    print(f'Beginning database import from slice {sliceMin} to {sliceMax}, '
          f'i.e. from PlaylistID {PlaylistID_Min} to {PlaylistID_Max}.')

    # Set the path of the MPD directory
    mpd_path = r'D:/Dropbox/IACS-CS-209-Spotify/mpd/data'
    # Move to this directory and get all filenames; each file is a slice
    os.chdir(mpd_path)
    fnames: List[str] = os.listdir()
    # The size of each slice
    slice_size: int = 1000

    # Make a sorted list of mpd slice files
    fnames_mpd: List[Tuple[str, int]] = list()
    # Iterate over all the files in this directory
    for fname in fnames:
        # Filenames have the format e.g. "mpd.slice.1000-1999.json"
        # First, check that this is an mpd data slice file; if not, skip it
        is_mpd_file: bool = fname.startswith("mpd.slice.") and fname.endswith(".json")
        if not is_mpd_file:
            continue
        # Extract the PlaylistID range from the file name
        pid_range: str = fname.split('.')[2]
        # The slice is the starting pid / 1000 (integer division)
        sliceNum: int = int(pid_range.split('-')[0]) // slice_size
        # If this slice is in the range, add (fname, sliceNum) to fnames_mpd
        if sliceMin <= sliceNum and sliceNum < sliceMax:
            fnames_mpd.append((fname, sliceNum))
    # Sort this by sliceNum
    fnames_mpd.sort(key=lambda x: x[1])

    # Get a database connection and a cursor; close the connection at the end!
    conn = getConnection()
    curs = conn.cursor()
    # Set mode to fast insertions on execute many
    curs.fast_executemany = True

    # Track progress
    num_processed: int = 0
    num_total: int = sliceMax - sliceMin
    # Start the timer
    t0 = time.time()

    for fname, sliceNum in fnames_mpd:
        # If we get here, this is a valid mpd data file in the range we want to process
        # Get the rows of data to insert into both tables
        rows_playlist, rows_playlist_entry = get_insertions(fname)
        # Range of PlaylistID in this slice
        PlaylistID_Min = min(pl[0] for pl in rows_playlist)
        PlaylistID_Max = max(pl[0] for pl in rows_playlist) + 1

```

```

# Delete records in this range on the table r.Playlist
delete_playlist(curs, PlaylistID_Min, PlaylistID_Max)
# Insert records in this range on the table r.Playlist
insert_playlist(curs, rows_playlist)

# Delete records in this range on the table r.PlaylistEntry
delete_playlist_entry(curs, PlaylistID_Min, PlaylistID_Max)
# Insert records in this range on the table r.PlaylistEntry
insert_playlist_entry(curs, rows_playlist_entry)

# Status update
num_processed += 1
num_left = num_total - num_processed
t1 = time.time()
elapsed_time = t1 - t0
average_pace = elapsed_time / num_processed
projected_time = num_left * average_pace
print(f'Processed slice {sliceNum} in {fname}. '
      f'Elapsed time {round(elapsed_time)}, projected time {round(projected_time)} seconds.')

# Close DB connection
curs.close()
conn.close()

```

Populating the Normalized DB Tables from the Raw Tables

Now that the two raw tables `r.Playlist` and `r.PlaylistEntry` are populated, it's straightforward to populate `Artist`, `Album`, `Track`, `Playlist` and `PlaylistEntry` in the `dbo` schema. I refer to this operation as a 'data import.' The scripts that do this are located in the repo folder `sql/03_DataImport`. Here are the ones that generate the five principal tables

Import Artist / 01_ImportArtist.sql

```

INSERT INTO dbo.Artist
(ArtistName, ArtistUri)
SELECT
    pe.ArtistName,
    pe.ArtistUri
FROM
    r.PlaylistEntry AS pe
GROUP BY
    pe.ArtistName, pe.ArtistUri
ORDER BY pe.ArtistUri;

```

Import Album / 02_ImportAlbum.sql

```

INSERT INTO dbo.Album
(AlbumName, AlbumUri)
SELECT
    pe.AlbumName,
    pe.AlbumUri
FROM
    r.PlaylistEntry AS pe
GROUP BY pe.AlbumName, pe.AlbumUri
ORDER BY pe.AlbumUri;

```

Import Track / 03_ImportTrack.sql


```

INSERT INTO dbo.Track
(TrackName, TrackUri, ArtistID, AlbumID)
SELECT
    pe.TrackName,
    pe.TrackUri,
    ar.ArtistID,
    al.AlbumID
FROM
    r.PlaylistEntry AS pe
    INNER JOIN dbo.Artist AS ar ON
        ar.ArtistUri = pe.ArtistUri
    INNER JOIN dbo.Album AS al ON
        al.AlbumUri = pe.AlbumUri
GROUP BY pe.TrackName, pe.TrackUri, ar.ArtistID, al.AlbumID
ORDER BY pe.TrackUri;

```

Import Playlist / 03_ImportPlaylist.sql

```

INSERT INTO dbo.Playlist
(PlaylistID, PlaylistName, NumTracks, NumArtists, NumFollowers, NumEdits,
DurationMS, IsCollaborative, ModifiedAt)
SELECT
    pl.PlaylistID,
    TRIM(pl.PlaylistName) AS PlaylistName,
    pl.NumTracks,
    pl.NumArtists,
    pl.NumFollowers,
    pl.NumEdits,
    pl.DurationMS,
    pl.IsCollaborative,
    pl.ModifiedAt
FROM
    r.Playlist AS pl;

```

Import PlaylistEntry / 04_ImportPlaylistEntry.sql

```

INSERT INTO dbo.PlaylistEntry
(PlaylistID, Position, TrackID)
SELECT
    pe.PlaylistID,
    pe.Position,
    t.TrackID
FROM
    r.PlaylistEntry AS pe
    INNER JOIN dbo.Track AS t ON
        t.TrackUri = pe.TrackUri;

```

Predicting the Last 10 Tracks of a Playlist

The Spotify app offers users a functionality suggesting tracks you might want to add to a playlist. It recommends 10 tracks. If you don't like any of them, it will recommend another 10, and so on. We set as our first task building a predictive model that could predict the last 10 tracks of a playlist by looking at features including the first $n - 10$ tracks and the name of the playlist. This problem is easy to state but hard to solve in practice. The data set is truly massive. There are 1,000,000 playlist and 2.26 million distinct tracks appearing. Scaling the problem down by considering only a smaller number of playlists at a time is a straightforward to make the scale more manageable. But we found it very challenging to deal with the huge number of columns.

The first idea we discussed was to simply cut off consideration of tracks that were not among the most popular, e.g. the top 100,000 tracks. There are a few issues with this. This data set has a very long tail. The first 100,000 tracks are dominant, but they only comprise 84% of the total number of entries of the on the PlaylistEntry table. (This calculation was done in the database view `v.TopTracks`). And even if you did reduce the size of this data set to 100,000 columns, the scale is still so massive that dense matrices are infeasible. With so many features, reducing the number of rows (playlists) you study is dangerous. Suppose for the sake of discussion we reduced it by a factor of 10, dropping down to 100,000. The matrix representing playlist entries would still have a size of 100,000 x 100,000, i.e. it would have 10^{10} or 10 billion entries in it. That is a lot of memory and is not going to be feasible.

If all we wanted to do was one or two basic linear algebra operations, then it is feasible load this data into sparse matrices. Care must be taken arrange the input into the right shape, and to only perform fast operations. As an example, if you want to perform matrix multiplication on $A * B$, this will work efficiently if A is in 'csr' (compressed sparse row) and B is in 'csc' (compressed sparse column) format. Otherwise it is going to be very slow.

The biggest problem with reducing the number of tracks under consideration though was the ambitious goal of the project: to produce recommendations that had a higher frequency of artists in a "promoted" group. In this case, the promoted group consisted of female mid-tier artists. Avriel generated a list of ~14,000 female midtier artists currently on Spotify, which matched up with about 10,000 artists in our dataset (out of about 295,000 artists total). These artists accounted for a tiny 0.8576% of total playlist recommendations with the default model which uses baseline frequencies. The upshot is that a model that limits the tracks to the top 100,000 or even the top 200,000 is going to be hopeless at recommending tracks deep in this long tail-it may as well be guessing.

As we contemplated the extent of this challenge and the limited time available, we chose to concentrate on three very straightforward models:

- Baseline Model: predict tracks based solely on their overall frequency in the training data set
- Playlist Name Model: predict tracks based on one feature, the name of the playlist
- Naive Bayes (Track Pair) Model: Predict a playlist by analyzing the frequency with which pairs of tracks appear together in the same playlist. Sum up the frequencies of the "other" tracks given the visible tracks in a playlist, and guess the top 10 other tracks If time permitted, we also considered adding a fourth model:
- Stacked Model: Predict playlist using a combination of the playlist name and track pair information.

When I considered these three models, I realized that none of the first three required any "fancy" machine learning and could all be implemented cleanly in SQL. Given the challenges of moving data from one platform to another and the limited time available, I chose to estimate the three basic models purely in SQL. The benefit of doing this in SQL is that is highly scalable and running on powerful hardware. All predictions were made on the full data set of one million playlists, split into 900,000 training and 100,000 test playlists.

Let's load in some additional dataframes so we can see examples of these predictions and their building blocks.

```
In [11]: f1 = ['Artist', 'Album', 'Track', 'Playlist', 'PlaylistEntry']
f2 = ['TrackRank', 'PlaylistName', 'PlaylistSimpleName']
frames = load_frames(f1 + f2)
```

```
Loaded Artist.
Loaded Album.
Loaded Track.
Loaded Playlist.
Loaded PlaylistEntry.
Loaded TrackRank.
Loaded PlaylistName.
Loaded PlaylistSimpleName.
```

```
Loaded 8 Data Frames.
Elapsed Time: 3.88 seconds.
```

The Baseline Model

This model is about as simple as you can get short of randomly guessing one of the 2.26 million tracks. Here is how it works in words. Assemble a list of the most popular tracks in descending order. The frequency here is how often this track appears as an entry in the training data set. To predict the last 10 tracks, take the list of candidates, remove the tracks that have already been played, and guess the top 10. That's it!

Below are two cells showing the SQL table definition for the persisted predictions in this model, and the SQL script used to generate the predictions. The exercise is conceptually trivial. All of the action comes in the handling of details like monitoring the progress and designing the insert batch so it can be restarted midway.

31_MakeTable_Prediction_Baseline.sql

```
CREATE TABLE dbo.Prediction_Baseline(  
  PlaylistID INT NOT NULL,  
  Position SMALLINT NOT NULL,  
  TrackID INT NOT NULL,  
  BaselineFrequency INT NOT NULL,  
  TrackRank INT NOT NULL,  
  CONSTRAINT PK_Prediction_Baseline_PlaylistID_Position  
    PRIMARY KEY (PlaylistID, Position),  
  CONSTRAINT UNQ_Prediction_Baseline_PlaylistID_TrackID  
    UNIQUE (PlaylistID, TrackID),  
  CONSTRAINT FK_Prediction_Baseline_PlaylistID  
    FOREIGN KEY (PlaylistID)  
      REFERENCES dbo.Playlist(PlaylistID),  
  CONSTRAINT FK_Prediction_Baseline_TrackID  
    FOREIGN KEY (TrackID)  
      REFERENCES dbo.Track(TrackID),  
);
```

01_Predict_Baseline.sql

```

DECLARE @PlaylistCount AS INT = 1000000;
DECLARE @BatchSize AS INT = 1000;
DECLARE @i AS INT = 0;
DECLARE @p1 INT;
DECLARE @p2 INT;

WHILE (@i * @BatchSize) < @PlaylistCount
BEGIN

    -- Range of playlists for this loop iteration
    SET @p1 = (@i * @BatchSize + 1);
    SET @p2 = ((@i+1) * @BatchSize);

    -- Delete records in this block
    DELETE pr
    FROM dbo.Prediction_Baseline AS pr
    WHERE pr.PlaylistID BETWEEN @p1 AND @p2;

    -- CTE to get (PlaylistID, TrackID, BaselineFrequency)
    WITH t1 AS (
    SELECT
        pl.PlaylistID,
        pl.NumTracks,
        tr.TrackID,
        tr.Frequency AS BaselineFrequency,
        tr.TrackRank AS TrackRank
    FROM
        -- Start with all playlists
        dbo.Playlist AS pl
        -- Track ranks up to 1024
        INNER JOIN dbo.TrackRank AS tr ON
            tr.TrackRank <= 1024
    WHERE
        -- Current block of PlaylistIDs
        pl.PlaylistID BETWEEN @p1 AND @p2
    ),
    t2 AS (
    SELECT
        t1.PlaylistID,
        row_number() OVER
            (PARTITION BY t1.PlaylistID ORDER BY t1.BaselineFrequency DESC)
            AS Position,
        t1.TrackID,
        t1.BaselineFrequency,
        t1.TrackRank
    FROM
        t1
    -- Don't pick tracks that are already on the first (n-10) elements of the playlist!
    WHERE NOT EXISTS
        (SELECT pe.PlaylistID FROM dbo.PlaylistEntry AS pe
        WHERE
            pe.PlaylistID = t1.PlaylistID AND
            pe.TrackID = t1.TrackID AND
            pe.Position <= t1.NumTracks - 11)
    )
    -- Insert the first 512 positions into the Prediction table
    INSERT INTO dbo.Prediction_Baseline
    (PlaylistID, Position, TrackID, BaselineFrequency, TrackRank)
    SELECT
        t2.PlaylistID,
        t2.Position,

```

```

t2.TrackID,
t2.BaselineFrequency,
t2.TrackRank
FROM
t2
WHERE
-- Only the first 512 positions
t2.Position <= 512;

-- Status update; manual Loop increment
PRINT CONCAT('Completed PlaylistID ', @i*@BatchSize+1, ' to ', (@i+1)*@BatchSize);
SET @i = @i+1;

END

```

Top 50 playlists in our Data Set

```
In [36]: # Start with the top 50 tracks
t1 = frames['TrackRank'].sort_values(by=['TrackRank']).head(50)
# Join the Track table to get the TrackName ArtistID
output_columns2 = ['TrackID', 'ArtistID', 'TrackName', 'TrackRank', 'Frequency']
t2 = pd.merge(left=t1, right=frames['Track'], on='TrackID', suffixes=('', '_dup'))
# Join the Artist table to get the ArtistName
output_columns3 = output_columns2 = ['TrackID', 'TrackName', 'ArtistName', 'TrackRank', 'Frequency']
t3 = pd.merge(left=t2, right=frames['Artist'], on='ArtistID', suffixes=('', '_dup'))[output_columns3]
display(t3.sort_values(by='TrackRank'))
```

	TrackID	TrackName	ArtistName	TrackRank	Frequency
0	2181619	HUMBLE.	Kendrick Lamar	1	41866
2	563659	One Dance	Drake	2	39143
6	2257296	Broccoli (feat. Lil Yachty)	DRAM	3	37164
7	2093570	Closer	The Chainsmokers	4	37015
10	920373	Congratulations	Post Malone	5	35941
13	1572888	Caroline	Aminé	6	31602
14	670143	iSpy (feat. Lil Yachty)	KYLE	7	31599
15	1309651	Bad and Boujee (feat. Lil Uzi Vert)	Migos	8	31475
16	2144121	XO TOUR Llif3	Lil Uzi Vert	9	31430
17	315491	Location	Khalid	10	31094
18	219481	Bounce Back	Big Sean	11	30378
20	1534119	Ignition - Remix	R. Kelly	12	29176
21	1756683	No Role Modelz	J. Cole	13	29035
22	247904	Mask Off	Future	14	28833
23	956975	I'm the One	DJ Khaled	15	28322
24	245916	No Problem (feat. Lil Wayne & 2 Chainz)	Chance The Rapper	16	28318
25	1848234	goosebumps	Travis Scott	17	28023
3	617161	Jumpman	Drake	18	28013
4	892549	Fake Love	Drake	19	27514
26	1526590	Despacito - Remix	Luis Fonsi	20	27464
8	1921832	Roses	The Chainsmokers	21	27281
9	204575	Don't Let Me Down	The Chainsmokers	22	27197
27	1719664	Gold Digger	Kanye West	23	27098
28	2217592	Shape of You	Ed Sheeran	24	26951
30	1022388	Redbone	Childish Gambino	25	26922
31	659132	Trap Queen	Fetty Wap	26	26363
11	1834762	White Iverson	Post Malone	27	26144
33	608401	The Hills	The Weeknd	28	26065
35	2255977	Riptide	Vance Joy	29	25646
36	1844621	Black Beatles	Rae Sremmurd	30	25493
34	1502294	Starboy	The Weeknd	31	25431
37	1636952	Panda	Designer	32	25416
38	729825	Ni**as In Paris	JAY Z	33	24844
39	2208504	Mr. Brightside	The Killers	34	24446
40	1420568	Lean On (feat. MØ & DJ Snake)	Major Lazer	35	24380
41	145476	That's What I Like	Bruno Mars	36	24240
42	524486	Needed Me	Rihanna	37	24064
1	1864886	DNA.	Kendrick Lamar	38	24007
43	1223838	Don't Stop Believin'	Journey	39	23785

	TrackID	TrackName	ArtistName	TrackRank	Frequency
12	2245093	rockstar	Post Malone	40	23483
44	1786395	Sorry	Justin Bieber	41	23443
45	1686992	1-800-273-8255	Logic	42	23391
19	1417502	I Don't Fuck With You	Big Sean	43	23368
5	678787	Hotline Bling	Drake	44	23278
46	2232627	Slide	Calvin Harris	45	23202
47	930260	Unforgettable	French Montana	46	22976
48	270971	Yeah!	Usher	47	22822
29	512923	Thinking Out Loud	Ed Sheeran	48	22615
32	1628340	679 (feat. Remy Boyz)	Fetty Wap	49	22555
49	1689780	Chill Bill	Rob \$tone	50	22262

The Playlist Simple Name Model

This model is the second simplest model we could think of. It only considers one feature of each playlist: the name of the playlist. While this might sound simple, it is also a powerful feature. The best way to see the power of this feature is to experiment with the Spotify app. When I was putting together playlists so I could complete our survey, I started by typing in the names of three playlists. I'm a classical music fan, and my playlists were titled "Virtuoso Piano", "Symphonies" and "Chamber Music". Based only on these titles, Spotify gave me a roster of entirely plausible suggestions, tilted to be sure to the most popular works in each field, but all relevant. The one million playlists included 73,498 distinct playlist names. A review of these playlist names shows there are a lot of effective duplicates separated only by letter casing and punctuation. Here is my favorite example: people like naming a playlist "vibes" with many variants...

```
In [43]: frames['PlaylistName'].query("PlaylistSimpleName == 'vibes'").head(20)
```

```
Out[43]:
```

	PlaylistNameID	PlaylistName	PlaylistSimpleNameID	PlaylistSimpleName
904	905	#Vibes	16196	vibes
905	906	#vibes	16196	vibes
906	907	#vibes#	16196	vibes
1118	1119	((((((Vibes))))))	16196	vibes
1144	1145	((Vibes))	16196	vibes
1579	1580	*^viBes	16196	vibes
1959	1960	*VIBES	16196	vibes
1960	1961	*Vibes*	16196	vibes
1961	1962	*vibes*	16196	vibes
2332	2333	.Vibes	16196	vibes
2472	2473	// vibes	16196	vibes
2473	2474	// VIBES //	16196	vibes
2627	2628	//Vibes//	16196	vibes
2714	2715	:(vibes	16196	vibes
2720	2721	: :) vibes	16196	vibes
2749	2750	:) Vibes	16196	vibes
2750	2751	:) vibes	16196	vibes
2795	2796	:: vibes ::	16196	vibes
2841	2842	:V:l:B:E:S:	16196	vibes
3335	3336	{vibes}	16196	vibes

For each playlist name, I computed a PlaylistSimpleName field by stripping out leading and trailing spaces (using the TRIM function), removing all punctuation, and moving all letters to lowercase. This reduced the count of distinct playlist simple names to

17,131. For each of these simple names, I then tabulated the frequency with which tracks appeared in all the playlists that shared that simple name. The prediction for a new playlist starts with candidates that are the most popular tracks sharing that playlist name. If necessary, this is padded by considering tracks with a high baseline frequency. (This is a rare corner case but it does happen if a playlist name is obscure and only appears a small number of times, perhaps once.) If a playlist name is not recognized, it is mapped to the empty string. This is also where playlist names containing only punctuation or non-ascii characters are mapped. Here is the SQL script that generates predictions using the playlist simple name model.

Build the rank (by frequency) of each track keyed by PlaylistSimpleNameID

12_Import_TrackRankBySimpleName.sql


```

TRUNCATE TABLE dbo.TrackRankBySimpleName;
GO

-- CTE with the core group by query computing frequency of each
-- (PlaylistSimpleNameID, TrackID) pair
WITH t1 AS (
SELECT
    pn.PlaylistSimpleNameID,
    pe.TrackID,
    COUNT(pe.TrackID) AS Frequency
FROM
    -- Start with all playlist entries
    dbo.PlaylistEntry AS pe
    -- The playlist for this entry
    INNER JOIN dbo.Playlist AS pl ON
        pl.PlaylistID = pe.PlaylistID
    -- Figure out if this playlist is train or test
    INNER JOIN dbo.TrainTestSplit AS tts ON
        tts.PlaylistID = pe.PlaylistID
    -- Get the PlaylistSimpleNameID from PlaylistName
    INNER JOIN dbo.PlaylistName AS pn ON
        pn.PlaylistName = pl.PlaylistName
WHERE
    -- Only training data!
    tts.TrainTestTypeID = 1
GROUP BY pn.PlaylistSimpleNameID, pe.TrackID
),
-- Join t1 against TrackRank to get BaselineFrequency and
-- compute the rank of each track sharing the SimpleName
t2 AS(
SELECT
    t1.PlaylistSimpleNameID,
    row_number() OVER
        (partition BY t1.PlaylistSimpleNameID
          ORDER BY t1.Frequency DESC, tr.Frequency DESC)
      AS TrackRank,
    t1.TrackID,
    t1.Frequency,
    tr.Frequency AS BaselineFrequency
FROM
    t1
    INNER JOIN dbo.TrackRank AS tr ON
        tr.TrackID = t1.TrackID
)
-- Insert into TrackRankBySimpleName from t2
INSERT INTO dbo.TrackRankBySimpleName
(PlaylistSimpleNameID, TrackRank, TrackID, Frequency, BaselineFrequency)
SELECT
    t2.PlaylistSimpleNameID,
    t2.TrackRank,
    t2.TrackID,
    t2.Frequency,
    t2.BaselineFrequency
FROM
    t2
WHERE
    -- Only compute ranks up to 1024 (need < 512 for prediction)
    t2.TrackRank <= 1024;

```

```

CREATE TABLE dbo.Prediction_SimpleName(
PlaylistID INT NOT NULL,
Position SMALLINT NOT NULL,
TrackID INT NOT NULL,
Frequency INT NOT NULL,
BaselineFrequency INT NOT NULL,
TrackRank INT NOT NULL,
CONSTRAINT PK_Prediction_SimpleName_PlaylistID_Position
    PRIMARY KEY (PlaylistID, Position),
CONSTRAINT UNQ_Prediction_SimpleName_PlaylistID_TrackID
    UNIQUE (PlaylistID, TrackID),
CONSTRAINT FK_Prediction_SimpleName_PlaylistID
    FOREIGN KEY (PlaylistID)
        REFERENCES dbo.Playlist(PlaylistID),
CONSTRAINT FK_Prediction_SimpleName_TrackID
    FOREIGN KEY (TrackID)
        REFERENCES dbo.Track(TrackID),
);

```

Predict tracks on a playlist by taking the top tracks sharing its PlaylistSimpleName, eliminating duplicates.

02_Predict_SimpleName.sql

```

USE SpotifyDB;
GO

DECLARE @PlaylistCount AS INT = 1000000;
DECLARE @BatchSize AS INT = 1000;
DECLARE @i AS INT = 0;
DECLARE @p1 INT;
DECLARE @p2 INT;

WHILE (@i * @BatchSize) < @PlaylistCount
BEGIN

    -- Range of playlists for this loop iteration
    SET @p1 = (@i * @BatchSize + 1);
    SET @p2 = ((@i+1) * @BatchSize);

    -- Delete records in this block
    DELETE pr
    FROM dbo.Prediction_SimpleName AS pr
    WHERE pr.PlaylistID BETWEEN @p1 AND @p2;

    -- t1 is Playlists joined with candidate tracks sharing the simple name of the playlist
    WITH t1 AS (
    SELECT
        pl.PlaylistID,
        pl.NumTracks,
        trsn.TrackRank,
        trsn.TrackID,
        trsn.Frequency,
        trsn.BaselineFrequency
    FROM
        -- Start with all playlists
        dbo.Playlist AS pl
        -- The simple name of the playlist
        INNER JOIN dbo.PlaylistName AS pn ON
            pn.PlaylistName = pl.PlaylistName
        -- The track frequencies for this playlist
        INNER JOIN dbo.TrackRankBySimpleName AS trsn ON
            trsn.PlaylistSimpleNameID = pn.PlaylistSimpleNameID
        -- Only the 1024 top candidates
        AND trsn.TrackRank <= 1024
    WHERE pl.PlaylistID BETWEEN @p1 AND @p2
    ),
    -- Filter out tracks that are already there; compute the Position
    t2 AS (
    SELECT
        t1.PlaylistID,
        row_number() OVER
            (PARTITION BY t1.PlaylistID ORDER BY t1.Frequency DESC, t1.BaselineFrequency DESC)
            AS Position,
        t1.TrackID,
        t1.Frequency,
        t1.BaselineFrequency,
        t1.TrackRank
    FROM
        t1
    -- Don't pick tracks that are already on the first (n-10) elements of the playlist!
    WHERE NOT EXISTS
        (SELECT pe.PlaylistID FROM dbo.PlaylistEntry AS pe
        WHERE
            pe.PlaylistID = t1.PlaylistID AND

```

```

        pe.TrackID = t1.TrackID AND
        pe.Position <= t1.NumTracks - 11)
    )
    -- Insert the first 512 positions into the Prediction table
    INSERT INTO dbo.Prediction_SimpleName
    (PlaylistID, Position, TrackID, Frequency, BaselineFrequency, TrackRank)
    SELECT
        t2.PlaylistID,
        t2.Position,
        t2.TrackID,
        t2.Frequency,
        t2.BaselineFrequency,
        t2.TrackRank
    FROM
        t2
    WHERE
        t2.Position <= 512;

    -- Status update; manual Loop increment
    PRINT CONCAT('Completed PlaylistID ', @i*@BatchSize+1, ' to ', (@i+1)*@BatchSize);
    SET @i = @i+1;

END

```

Brief Digression: Behind the Decision to Predict in SQL

My original plan was to organize the back end of the data in SQL server, then export all the normalized data in frames, and do all predictions and analysis in Python. I had grand dreams of training classifiers such as logistic regression using sparse matrices as inputs, then graduating to training more sophisticated neural networks that could identify subtle features. These dreams crashed into two harsh realities. The first was my lack of familiarity with Pandas for doing advanced querying, indexing, and joining. I made some progress overcoming this issue, but the second challenge was the real deal breaker. The initial model I wrote to implement the simple name classification model above ran far too slowly to be practical. After extensive efforts to speed it up, the tqdm progress bar estimated that this model would take 106 hours to predict our entire data set. These predictions would have then needed to be persisted on disk so they could be accessed again in the future. It also would have been a fair amount of additional work to find a good way to stage this calculation into smaller independent chunks. In contrast to this, the SQL queries above predicted the next 10 tracks on all one million playlists in a runtime of approximately one hour. (58 minutes to be more precise). So it ran over 100 times faster. Of course, I am not claiming that it would have been impossible to get a faster runtime on a Python program. My point is that after I made a solid effort to write a program that was not slow, it was still way slower than using the SQL Server back end. That was when I decided to make the predictions in our simple models using SQL queries.

The Naive Bayes Model Using Track Pairs

The Naive Bayes model is the first "nontrivial" model we came up with. Here is a brief overview of Naive Bayes classifiers on Wikipedia: https://en.wikipedia.org/wiki/Naive_Bayes_classifier (https://en.wikipedia.org/wiki/Naive_Bayes_classifier). For a more in-depth discussion on Naive Bayes classifiers, please see Pattern Recognition and Machine Learning by Chris Bishop (2006), page 380. In this case, the underlying idea is that we can model the probability that one track T_j appears in a playlist conditional on the presence of another track T_i . We can assemble the frequency with which any pair of tracks appears on the same playlist. The intuition behind this is that if two tracks appear together frequently, those tracks are "similar" in the sense that people who put one of them on a playlist are likely to put a second one on a playlist. As an example, my wife had as one of her three playlists on the survey a "kids" playlist that included soundtracks from Disney movies. The Naive Bayes classifier predicted that a playlist including "Hakuna Matata" from "The Lion King" was likely to also include the track "Under the Sea" from "The Little Mermaid". The frequency with which a pair of tracks appears together reflects both how correlated two tracks are as well as how broadly popular they are.

In order to make predictions using the Naive Bayes Track Pair model, I first generated a table called TrackPair. This table counted how often each pair of tracks (Track_i, Track_j) appeared together in the training data set. This is a big table; it contains 2.12 billion records. On the other hand, it is still massively smaller than the number of entries in a dense matrix pairing up tracks. Since there are 2.26 million tracks, there are 5.11 trillion entries in a dense matrix of track pairs. The "sparsity factor" of this matrix is thus 1 in 2410 entries are non-zero. Once all of the track pairs have been tabulated, here is how the model would predict the next

10 tracks on a playlist. For each of the visible tracks T_i in the playlist, it would add up the frequency of all the tracks T_j that appeared together with T_i on the same playlist. In terms of the matrix of track pairs, it would take row i . These would be added up together to generate a list of empirical frequencies. After removing duplicates, the model would predict the 10 tracks with the highest frequency.

These models can be described very simply in terms of the matrix of playlist entries. Let PE be an n_p by n_t matrix, where n_p is the number of playlists, n_t is the number of tracks, and $PE_{i,j}$ is a 1 when playlist i contains track j and a zero otherwise. Then the $n_t \times n_t$ matrix of track pairs is given by

$$TP = X^T \cdot X$$

Now suppose that p is an $n_t \times 1$ column vector representing the visible tracks in a playlist. The matrix product

$$y = TP \cdot p$$

is an $n_t \times 1$ column vector of predicted frequencies. This can be generalized to a matrix product

$$Y = TP \cdot X^T$$

The i th column of this matrix would have the frequencies used to predict playlist i .

A more convenient matrix product for prediction would thus be

$$Y = X \cdot TP$$

(here using the fact that $TP = TP^T$ because TP is a symmetric matrix). The matrix Y defined above has size $n_p \times n_t$. By searching for the largest 10 entries in row i , we can predict the next 10 tracks for playlist i .

I experimented with this enough to see that it would in fact have been a viable approach for making predictions in terms of performance so long as all the matrices were sparse and with optimal representations. By this point however I was already committed to predicting the Track Pair model in SQL and there was no benefit to switching back to Python sparse matrix multiplications to do it.

Limiting the Band Width of the Track Pair Model

Once the TrackPair table has been generated in the database, it is very simple to write a query corresponding to the model described above. Unfortunately when I tried writing this query, it was on track to take days to complete. I first tried to improve the performance of by creating a helper table called PlaylistTrack_Visible that had primary key (PlaylistID, TrackID). This did indeed help, because the original PlaylistEntry table is keyed by (PlaylistID, Position) and it only had an index on (PlaylistID, TrackID). Still it wasn't enough. There were too many entries that had to be summed in the group by clause of the query. The idea that made it feasible to make predictions on this model was to take a **band limited** version of the track pair model.

Recall that when taking the contributions of a visible track on a playlist to the predicted tracks, the Naive Bayes model is adding up the nonzero terms in row i of the TrackPair matrix. The idea of the band limited matrix is simple and obvious: consider only the 64 largest entries in each row of the TrackPair matrix. The database table TrackPair_BL64 does exactly that. It also employs a second strategy to limit the size of the matrix. It only considers entries that are at least 3. Pairs of tracks that appeared only once or twice are eliminated from consideration. These two adjustments shrink down the TrackPair matrix by a very large factor: from 2.12 billion entries to just 24.6 million entries, a factor of 86. The information content is not meaningfully degraded. Predicting with this more manageably sized bandwidth limited track pair table allowed all one million playlists to be predicted in about three and a half hours. Below I have included the highlights of the SQL implementation of predictions in this model.

One might also point out that since the Track Pair matrix is symmetric, the size of these tables could be cut down in half by only storing entries where $i \leq j$. This is absolutely true. On the other hand, doing this would require joining two copies of the table. Extensive experiments over the years with exactly this issue (storing symmetric tables used to join in queries) have convinced me that by spending more space to store the full symmetric version of the table, you can often gain a large amount of time. I have plenty of space on my hard drives (a RAID of SSDs with a total storage capacity of 15 terabytes), but very limited time to complete this assignment. So I opted to trade space to gain time.

Build & Populate the TrackPair Table: All 2.12 Billion Records!

17_MakeTable_TrackPair.sql

```

CREATE TABLE dbo.TrackPair(
TrackID_1 INT NOT NULL,
TrackID_2 INT NOT NULL,
Frequency INT NOT NULL,
-- Primary key is the pair (TrackID_1, TrackID_2)
CONSTRAINT PK_TrackPair_TrackID_1_TrackID_2 PRIMARY KEY (TrackID_1, TrackID_2),
-- Foreign keys onto Track table
CONSTRAINT FK_TrackPair_TrackID_1
FOREIGN KEY (TrackID_1) REFERENCES dbo.Track(TrackID),
CONSTRAINT FK_TrackPair_TrackID_2
FOREIGN KEY (TrackID_2) REFERENCES dbo.Track(TrackID),
);

```

09_Import_TrackPair.sql

```

INSERT INTO dbo.TrackPair
(TrackID_1, TrackID_2)
SELECT
    pe1.TrackID AS TrackID_1,
    pe2.TrackID AS TrackID_2,
    COUNT(pl.PlaylistID) AS Frequency
FROM
    -- Start with playlists
    dbo.Playlist AS pl
    -- Train, Val or Test?
    INNER JOIN dbo.TrainTestSplit AS tts ON
        tts.PlaylistID = pl.PlaylistID
    -- Get two copies of all the tracks on this playlist
    INNER JOIN dbo.PlaylistEntry AS pe1 ON
        pe1.PlaylistID = pl.PlaylistID
    INNER JOIN dbo.PlaylistEntry AS pe2 ON
        pe2.PlaylistID = pl.PlaylistID
WHERE
    -- Only the training set!
    tts.TrainTestTypeID = 1
GROUP BY
    pe1.TrackID, pe2.TrackID

```

Build and Populate the TrackPair_BL64 Table from TrackPair: Now Only 24.6 Million Records

17_MakeTable_TrackPair_BL64.sql

```

CREATE TABLE dbo.TrackPair_BL64(
TrackID_1 INT NOT NULL,
TrackID_2 INT NOT NULL,
Frequency INT NOT NULL,
-- Primary key is the pair (TrackID_1, TrackID_2)
CONSTRAINT PK_TrackPair_BL64 PRIMARY KEY (TrackID_1, TrackID_2)
);

```

09_Import_TrackPair_BL64.sql

```

DECLARE @TrackCount AS INT;
DECLARE @BatchSize AS INT = 10000;
DECLARE @i AS INT = 0;
DECLARE @t1 INT;
DECLARE @t2 INT;

-- Get Last TrackID
SELECT
    @TrackCount = COALESCE(MAX(tr.TrackID),0)
FROM
    dbo.Track AS tr;

WHILE (@i * @BatchSize) < @TrackCount
BEGIN

    -- Range of playlists for this loop iteration
    SET @t1 = (@i * @BatchSize + 1);
    SET @t2 = ((@i+1) * @BatchSize);

    WITH t1 AS (
    SELECT
        tp.TrackID_1,
        tp.TrackID_2,
        tp.Frequency,
        row_number() OVER
            (partition BY tp.TrackID_1 ORDER BY Frequency DESC)
            AS ColumnRank
    FROM
        dbo.TrackPair AS tp
    WHERE
        -- Selected range of TrackIDs
        tp.TrackID_1 BETWEEN @t1 AND @t2 AND
        -- Only use entries with frequency > 2
        tp.Frequency > 2 AND
        -- Don't return TrackID_1 again (obviously)
        tp.TrackID_2 <> tp.TrackID_1
    )
    INSERT INTO dbo.TrackPair_BL64
    (TrackID_1, TrackID_2, Frequency)
    SELECT
        t1.TrackID_1,
        t1.TrackID_2,
        t1.Frequency
    FROM
        t1
    WHERE
        -- Only the top 64 tracks arising from the TrackID_1
        t1.ColumnRank <= 64;

    -- Status update; manual loop increment
    PRINT CONCAT('Completed TrackID ', @t1, ' to ', @t2);
    SET @i = @i+1;

END

```

With the TrackPair_BL64 table available, the strategy for predicting in the Naive Bayes Track Pair model is simple. Start with the "visible" playlist entries (those in the first \$n-10\$ positions), then join each visible track onto the TrackPair_BL64 table, and group by two fields: the PlaylistID being predicted and the predicted track \$T_j\$. Here is the SQL that does this.

03_Predict_TrackPair.sql


```

DECLARE @PlaylistCount AS INT = 1000000;
DECLARE @BatchSize AS INT = 1000;
DECLARE @i AS INT = 0;
DECLARE @p1 INT;
DECLARE @p2 INT;

WHILE (@i * @BatchSize) < @PlaylistCount
BEGIN

    -- Range of playlists for this loop iteration
    SET @p1 = (@i * @BatchSize + 1);
    SET @p2 = ((@i+1) * @BatchSize);

    -- Delete records in this block
    DELETE pr
    FROM dbo.Prediction_TrackPair AS pr
    WHERE pr.PlaylistID BETWEEN @p1 AND @p2;

    WITH t1 AS(
    SELECT
        ptv.PlaylistID,
        tp.TrackID_2 AS TrackID,
        SUM(tp.Frequency) AS Frequency
    FROM
        -- Start with visible playlist entries
        dbo.PlaylistTrack_Visible AS ptv
        -- For each track on the visible portion, find all its pairs
        -- Use the bandwidth limited version (top 64 columns per row)
        -- to speed performance
        INNER JOIN dbo.TrackPair_BL64 AS tp ON
            tp.TrackID_1 = ptv.TrackID
    WHERE
        ptv.PlaylistID BETWEEN @p1 AND @p2
    GROUP BY ptv.PlaylistID, tp.TrackID_2
    ),
    -- Get the track rank of each candidate
    t2 AS(
    SELECT
        t1.PlaylistID,
        t1.TrackID,
        row_number() OVER
            (partition BY t1.PlaylistID ORDER BY t1.Frequency DESC)
            AS TrackRank,
        t1.Frequency
    FROM
        t1
    ), t3 AS (
    -- Limit these to the 1024 most common tracks and
    -- filter out tracks already there
    SELECT
        t2.PlaylistID,
        row_number() OVER
            (partition BY t2.PlaylistID ORDER BY t2.Frequency DESC, trr.Frequency DESC)
            AS Position,
        t2.TrackID,
        t2.Frequency,
        trr.Frequency AS BaselineFrequency,
        t2.TrackRank
    FROM
        t2
        -- Join Playlist to get the number of tracks

```

```

INNER JOIN dbo.Playlist AS p1 ON
    p1.PlaylistID = t2.PlaylistID
-- Join TrackRank to get the baseline frequency
INNER JOIN dbo.TrackRank AS trr ON
    trr.TrackID = t2.TrackID
WHERE
    -- Only want at most 512 tracks per playlist here
    t2.TrackRank <= 512 AND
    -- No duplicates
    NOT EXISTS (
        SELECT ptv2.PlaylistID FROM dbo.PlaylistTrack_Visible AS ptv2
        WHERE ptv2.PlaylistID = t2.PlaylistID AND ptv2.TrackID = t2.TrackID
    )
)
INSERT INTO dbo.Prediction_TrackPair
(PlaylistID, Position, TrackID, Frequency, BaselineFrequency, TrackRank)
SELECT
    t3.PlaylistID,
    t3.Position,
    t3.TrackID,
    t3.Frequency,
    t3.BaselineFrequency,
    t3.TrackRank
FROM t3
WHERE
    t3.Position <= 256;

-- Status update; manual Loop increment
PRINT CONCAT('Completed PlaylistID ', @p1, ' to ', @p2);
SET @i = @i+1;

END

```

Scoring Predictions in the Models

Once a model has made its predictions on the training and test set, we score it by comparing the top 100 tracks it predicts against the actual last 10 tracks of each playlist in both train and test sets. We disregarded the position (order) of the last 10 tracks to simplify the analysis. We did account for the position we predicted, because it's a lot better to make a correct guess in position 1 than in position 100. We computed the average number of correct guesses in each position from 1 to 100, then added them up to form a cumulative score we could visualize.

The SQL to do these calculations is simple. First an intermediate table called `PlaylistEntry_Last10` is built up. This table is keyed by (PlaylistID, TrackID) and contains one record for each of the ten tracks that were at the end of a given playlist. The advantage of this auxiliary table as compared to the original one is that it is smaller (it has only the hits) and is keyed with TrackID rather than position, making the query run much faster. In fact, the scoring query is so efficient that the score calculation is not persisted in a table at all, but is just a database view.

Generate the `PlaylistEntry_Last10` Table: These are the Tracks That are "Hits" When We Predict Them

18_Import_PlaylistEntry_Last10.sql

```

DECLARE @PlaylistCount AS INT = 1000000;
DECLARE @BatchSize AS INT = 1000;
DECLARE @i AS INT = 0;
DECLARE @p1 INT;
DECLARE @p2 INT;

WHILE (@i * @BatchSize) < @PlaylistCount
BEGIN

    -- Range of playlists for this loop iteration
    SET @p1 = (@i * @BatchSize + 1);
    SET @p2 = ((@i+1) * @BatchSize);

    INSERT INTO dbo.PlaylistTrack_Last10
    (PlaylistID, TrackID)
    SELECT DISTINCT
        pe.PlaylistID,
        pe.TrackID
    FROM
        dbo.Playlist AS pl
        INNER JOIN dbo.PlaylistEntry AS pe ON
            pl.PlaylistID = pe.PlaylistID AND
            pe.Position > pl.NumTracks - 11
    WHERE
        pl.PlaylistID BETWEEN @p1 AND @p2

    -- Status update; manual loop increment
    PRINT CONCAT('Completed PlaylistID ', @i*@BatchSize+1, ' to ', (@i+1)*@BatchSize);
    SET @i = @i+1;

END

```

Score the TrackPair Model on Train and Test

11_Scores_TrackPair.sql

```

CREATE VIEW v.Scores_TrackPair AS
-- Hit rate on Training set
WITH trn AS (
SELECT
    pr.Position,
    COUNT(pr.TrackID) AS Hits,
    COUNT(pr.TrackID) / 900000.0 AS HitRate
FROM
    dbo.Prediction_TrackPair AS pr
    INNER JOIN dbo.PlaylistTrack_Last10 AS ptl ON
        ptl.PlaylistID = pr.PlaylistID AND
        ptl.TrackID = pr.TrackID
    -- Only training data
    INNER JOIN dbo.TrainTestSplit AS tts ON
        tts.PlaylistID = pr.PlaylistID AND
        tts.TrainTestTypeID = 1
WHERE
    -- Only score the top 100 guesses
    pr.Position <= 100
GROUP BY pr.Position
),
-- Hit rate on Test set
tst AS (
SELECT
    pr.Position,
    COUNT(pr.TrackID) AS Hits,
    COUNT(pr.TrackID) / 100000.0 AS HitRate
FROM
    dbo.Prediction_TrackPair AS pr
    INNER JOIN dbo.PlaylistTrack_Last10 AS ptl ON
        ptl.PlaylistID = pr.PlaylistID AND
        ptl.TrackID = pr.TrackID
    -- Only training data
    INNER JOIN dbo.TrainTestSplit AS tts ON
        tts.PlaylistID = pr.PlaylistID AND
        tts.TrainTestTypeID = 3
WHERE
    -- Only score the top 100 guesses
    pr.Position <= 100
GROUP BY pr.Position
)
SELECT
    trn.Position,
    trn.HitRate AS HitRate_Trn,
    tst.HitRate AS HitRate_Tst,
    SUM(trn.HitRate) OVER (ORDER BY trn.Position) AS CumHits_Trn,
    SUM(tst.HitRate) OVER (ORDER BY tst.Position) AS CumHits_Tst
FROM
    trn
    INNER JOIN tst ON
        tst.Position = trn.Position;

```

Let's take a quick look at the scores for the TrackPair model:

```
In [47]: frames = load_frames(['Scores_TrackPair'])
```

Loaded Scores_TrackPair.

Loaded 1 Data Frames.

Elapsed Time: 0.01 seconds.

```
In [48]: frames['Scores_TrackPair']
```

```
Out[48]:
```

	Position	HitRate_Trn	HitRate_Tst	CumHits_Trn	CumHits_Tst
0	1	0.067654	0.06520	0.067654	0.06520
1	2	0.055394	0.05148	0.123049	0.11668
2	3	0.048843	0.04626	0.171892	0.16294
3	4	0.045328	0.04252	0.217220	0.20546
4	5	0.041890	0.04164	0.259110	0.24710
5	6	0.039272	0.03778	0.298382	0.28488
6	7	0.036716	0.03564	0.335098	0.32052
7	8	0.035367	0.03405	0.370464	0.35457
8	9	0.032817	0.03272	0.403281	0.38729
9	10	0.031576	0.03094	0.434857	0.41823
10	11	0.030004	0.03043	0.464861	0.44866
11	12	0.028313	0.02734	0.493174	0.47600
12	13	0.027561	0.02673	0.520736	0.50273
13	14	0.026560	0.02646	0.547296	0.52919
14	15	0.025948	0.02505	0.573243	0.55424
15	16	0.025112	0.02424	0.598355	0.57848
16	17	0.024464	0.02378	0.622820	0.60226
17	18	0.023833	0.02330	0.646653	0.62556
18	19	0.022860	0.02401	0.669513	0.64957
19	20	0.022353	0.02199	0.691867	0.67156
20	21	0.021860	0.02090	0.713727	0.69246
21	22	0.021318	0.02111	0.735044	0.71357
22	23	0.020823	0.02087	0.755868	0.73444
23	24	0.020631	0.02063	0.776499	0.75507
24	25	0.020508	0.01986	0.797007	0.77493
25	26	0.019760	0.01966	0.816767	0.79459
26	27	0.019488	0.01920	0.836254	0.81379
27	28	0.019432	0.01920	0.855687	0.83299
28	29	0.018969	0.01907	0.874655	0.85206
29	30	0.018844	0.01867	0.893500	0.87073
...
70	71	0.013607	0.01282	1.518726	1.48831
71	72	0.013514	0.01310	1.532241	1.50141
72	73	0.013084	0.01291	1.545325	1.51432
73	74	0.013016	0.01228	1.558341	1.52660
74	75	0.012939	0.01236	1.571280	1.53896
75	76	0.012822	0.01205	1.584102	1.55101
76	77	0.012442	0.01212	1.596544	1.56313
77	78	0.012393	0.01169	1.608937	1.57482
78	79	0.012217	0.01210	1.621154	1.58692
79	80	0.012309	0.01181	1.633463	1.59873
80	81	0.012001	0.01193	1.645464	1.61066
81	82	0.011986	0.01186	1.657450	1.62252

	Position	HitRate_Trn	HitRate_Tst	CumHits_Trn	CumHits_Tst
82	83	0.011833	0.01173	1.669283	1.63425
83	84	0.011921	0.01173	1.681204	1.64598
84	85	0.011956	0.01091	1.693160	1.65689
85	86	0.011663	0.01181	1.704823	1.66870
86	87	0.011479	0.01091	1.716302	1.67961
87	88	0.011492	0.01140	1.727794	1.69101
88	89	0.011474	0.01102	1.739269	1.70203
89	90	0.011284	0.01119	1.750553	1.71322
90	91	0.011149	0.01118	1.761702	1.72440
91	92	0.010898	0.01069	1.772600	1.73509
92	93	0.011054	0.01159	1.783654	1.74668
93	94	0.011008	0.01060	1.794662	1.75728
94	95	0.010834	0.01055	1.805496	1.76783
95	96	0.010883	0.01036	1.816380	1.77819
96	97	0.010683	0.01058	1.827063	1.78877
97	98	0.010750	0.01039	1.837813	1.79916
98	99	0.010306	0.01063	1.848118	1.80979
99	100	0.010549	0.01006	1.858667	1.81985

100 rows × 5 columns

Predicting With the Stacked Model

The idea of the Stacked Model was to create an ensemble prediction combining the Simple Name and Track Pair models. This idea seemed promising and I had high hopes for it. Due to very limited time when I attempted to estimate this model, I was unable to get both underlying predictors into Python to try to train a "real" model like a logistic regression using the predicted frequencies for tracks between playlist name and track pair. Instead, I took a quick and dirty estimate of this model in SQL. I assembled the top 256 candidates in each component model and converted the frequency of each candidate track into a pseudo-probability by dividing by the total frequency of that predictor. If a track appeared in one category (e.g. simple name) but not the other (e.g. track pair), its pseudo-probability in the track pair model was recorded as zero. This operation in SQL was accomplished with a combination of the UNION clause and a coalesce function that mapped missing (null) frequencies to zero.

Through trial and error, I found that a combination of 75% on the TrackPair model and 25% on the Simple Name model appeared to do reasonably well on a subset of the test set (the first 10,000 rows out of 100,000). I estimated this whole model and was disappointed to see that it performed marginally worse than the original TrackPair model. I am not ready to give up entirely on the idea of ensembling predictions based on both the playlist name and track pair information. But for the available time on this project it did not work as well as the more straightforward Track Pair model, so I will not present the details of its implementation here.

Idea for the Future: k-NN Using a Playlist Pair Metric

We've seen above how to construct a Track Pair matrix

$$TP = X^T \cdot X$$

We could apply the exact same idea, but use matrix transposes to generate a Playlist Pair matrix as follows:

$$PP = X \cdot X^T$$

Just as the Track Pair matrix is an indication of the strength of correlation between two tracks, the playlist pair matrix will have high entries on cells (i, j) where the two playlists are similar. If we converted from the raw frequencies to a correlation by dividing out by the number of entries in the two playlist, we could attempt create a notion of correlation between two playlists. We could then attempt to predict the next tracks on a playlists by finding its nearest neighbors.

I believe that this model is quite promising. Unfortunately I did not have time to pursue it. I attempted to estimate it but the calculation bogged down and I did not have enough time to bring it to completion. I spent enough time working on it though to identify it as a promising avenue of exploration in the future if I came back to this problem.

Recommending Playlists to Survey Participants

As part of our experimental design, we generated recommendations for all survey participants using two strategies. The first strategy was the Naive Bayes model using track pairs. This was a "neutral" strategy and was selected because it had the best performance in predicting the last 10 tracks in the test set. The second strategy was a modified track pair model. The modification was designed to increase the rate at which a selected group of promoted artists had their tracks recommended. In this case, the promoted artists comprised a sample of all the female mid-tier artists on the Spotify system as of the time that Avriel ran a query to that effect. Her query generated a list of approximately 14,000 artists, keyed by Spotify's ArtistUri field. This table was imported into the database table named PromotedArtist. We were only able to recommend artists who were present in our sample data; there was an overlap of 10,082 artists.

The neutral recommendation strategy was essentially identical to the track pairs strategy in the training and test sets. The only practical implementation difference is that all of the tracks in the survey responder's playlist were visible; in the training and testing we needed to hold out the last 10. Below is an example of the layout of the survey data. To protect user privacy, I will only share my own survey responses.

```
In [134]: fl = ['SurveyResponse', 'SurveyPlaylist', 'SurveyPlaylistEntry', 'Track',  
              'SurveyRecommendations', 'SurveyRecommendationsPromoted']  
frames = load_frames(fl)
```

```
Loaded SurveyResponse.  
Loaded SurveyPlaylist.  
Loaded SurveyPlaylistEntry.  
Loaded Track.  
Loaded SurveyRecommendations.  
Loaded SurveyRecommendationsPromoted.
```

```
Loaded 6 Data Frames.  
Elapsed Time: 1.93 seconds.
```

One Survey Response Record

```
In [135]: # Filter the SurveyResponse frame for one user's response  
df = frames['SurveyResponse']  
idx = df.index[df.RecipientEmail == 'michael.s.emanuel@gmail.com']  
df.iloc[idx.values[0]]
```

```
Out[135]: SurveyResponseID                20  
RecipientName                Michael S. Emanuel  
RecipientEmail              michael.s.emanuel@gmail.com  
SpotifyUserName              1262176408  
PlaylistUrl_1      https://open.spotify.com/user/1262176408/playl... (https://open.spotify.com/user/12  
62176408/playl...)  
PlaylistUrl_2      https://open.spotify.com/user/1262176408/playl... (https://open.spotify.com/user/12  
62176408/playl...)  
PlaylistUrl_3      https://open.spotify.com/user/1262176408/playl... (https://open.spotify.com/user/12  
62176408/playl...)  
PlaylistUri_1      6ypozS07dBAM72MGVe7zkH?si=2w4b79U2TySRiGhrY-8EKA  
PlaylistUri_2      60fNcluTJJgs3iEC3ekOZb?si=Hn1CPJqqT7uPk19ahBzmjg  
PlaylistUri_3      2gN8PrwCh1NXTSjiXaqoOw?si=-OdePYScTuCNNTN7z5z0IA  
Name: 19, dtype: object
```

Three Playlists Submitted by One Survey Participant

```
In [136]: # The SurveyResponseID of this survey participant
sr_id = frames['SurveyResponse'].iloc[idx].SurveyResponseID.values[0]
# A mask to filter just this user's responses back
mask_pl = frames['SurveyPlaylist'].SurveyResponseID == sr_id
frames['SurveyPlaylist'][mask_pl]
```

Out[136]:

	SurveyPlaylistID	SurveyResponseID	PlaylistNum	PlaylistUri	PlaylistName	PlaylistSimpleName
57	58	20	1	6ypozSO7dBAM72MGVe7zkH? si=2w4b79U2TySRiGHRy-8EKA	Virtuoso Piano	NaN
58	59	20	2	60fNcluTJJgs3iEC3ekOZb? si=Hn1CPJqqT7uPkl9ahBzmjg	Symphonies	NaN
59	60	20	3	2gN8PrwCh1NXTSjiXaqOw?si=- OdePYScTuCNNTN7z5zOIA	Chamber Music	NaN

Above we can see the three playlists I submitted. The names of my playlists were "Virtuoso Piano", "Symphonies" and "Chamber Music". The three NaN entries are a jarring reality check about the limitations of the simple playlist name approach implemented here. Even over one million playlists, all three of the above playlist titles were unrecognized! (I double checked this with a manual query on the database, and it's not a coding error. There are playlists named "Symphonic", "symphonic", and "symphonic!!!" in the MPD, but nobody named a playlist "Symphonies"). These examples are of course very challenging because Spotify is primarily a popular music service and classical music fans are a quite small constituency. Let's explore the tracks on the Virtuoso Piano playlist to see the inputs considered by the Track Pairs model.

Playlist Entries on One Participant's Playlists


```
In [137]: # The three SurveyPlaylistIDs
pids = frames['SurveyPlaylist'][mask_pl][['SurveyPlaylistID']]
# The playlist entries (IDs only)
df = frames['SurveyPlaylistEntry']
# Filter the survey playlist entries to the three selected playlists
df = pd.merge(left=pids, right=df, on='SurveyPlaylistID')
# Join these entries on the track table
output_cols = ['SurveyPlaylistID', 'Position', 'TrackName']
df = pd.merge(left=df, right=frames['Track'], on='TrackID')[output_cols]
df
```

Out[137]:

	SurveyPlaylistID	Position	TrackName
0	58	24	Pictures at an Exhibition: Promenade
1	58	25	Pictures at an Exhibition: Gnomus
2	58	26	Pictures at an Exhibition: Promenade
3	58	27	Pictures at an Exhibition: Il vecchio castello
4	58	28	Pictures at an Exhibition: Promenade
5	58	29	Pictures at an Exhibition: Tuileries
6	58	30	Pictures at an Exhibition: Bydlo
7	58	31	Pictures at an Exhibition: Promenade
8	58	32	Pictures at an Exhibition: Ballet of the Unhat...
9	58	33	Pictures at an Exhibition: Samuel Goldenberg &...
10	58	35	Pictures at an Exhibition: Limoges marché
11	58	36	Pictures at an Exhibition: Catacombae (Sepulcr...
12	58	37	Pictures at an Exhibition: Con mortuis in ling...
13	58	38	Pictures at an Exhibition: The Hut on Fowl's L...
14	58	39	Pictures at an Exhibition: The Great Gate of Kiev
15	59	5	Mozart: Symphony No. 40 in G Minor, K. 550: I....
16	59	6	Mozart: Symphony No. 40 in G Minor, K. 550: II...
17	59	7	Symphony No. 40 in G minor K550 [version with ...
18	59	8	Mozart: Symphony No. 40 in G Minor, K. 550 [ve...
19	59	21	Symphony No.8 In F, Op.93: 1. Allegro vivace e...
20	59	22	Symphony No.8 In F, Op.93: 2. Allegretto scher...
21	59	23	Symphony No.8 In F, Op.93: 3. Tempo di menuetto
22	59	24	Symphony No.8 In F, Op.93: 4. Allegro vivace
23	60	5	Clarinet Quintet in A Major: Clarinet Quintet ...

This table shows the information content available to Track Pair prediction model. If you had seen my original playlists and had a passing familiarity with classical music, you would see that the prediction task ranges from quite challenging to almost hopeless due solely to poor data coverage. The first playlist I submitted, Virtuoso Piano, had 46 tracks. It included Beethoven piano sonatas, Chopin etudes, Schubert impromptus, and Pictures at an Exhibition by Moussorgsy. Of these 46 tracks, only the 15 tracks in Pictures at an Exhibition were recognized by the data set. The Symphonies playlist fared about as badly. 24 tracks dropped down to 8. The Chamber Music playlist was almost a complete whiff. Out of 19 tracks including some of the most popular chamber music works in the repertory, only one track (a single movement of Mozart's A major clarinet quintet) was recognized. This example is a self contained demonstration of why it is so hard to limit the number of tracks and / or playlists under consideration in this problem if you want to do anything other than make plausible recommendations to people who are interested in the most popular tracks. The world of music is large and deep.

The mechanical steps of generating survey recommendations in the unadjusted Track Pair model are so similar the code presented above that I will omit it here. The only interesting detail is that I used the full TrackPair table rather than the bandwidth limited table. This was because some survey participants weren't getting the full slate of 10 recommendations if the tracks on their initial playlist had only a small number of hits. A more interesting question is how we generated recommendations that promoted female midtier artists. This is described below.

Recommendations to Extend the "Virtuoso Piano" Playlist

```
In [141]: # The survey playlist ID to explore
spid = 58
df = frames['SurveyRecommendations']
mask = (df.SurveyPlaylistID == spid)
output_cols = ['Position', 'Frequency', 'TrackName']
df[mask][output_cols]
```

Out[141]:

	Position	Frequency	TrackName
570	1	17	Toccata, Adagio & Fugue in C Major, BWV 564: P...
571	2	16	Adagio for Strings from the String Quartet, Op...
572	3	16	Adagio for Strings, Op.11
573	4	16	Pines Of Rome
574	5	16	Symphony No.5 In C Sharp Minor: 4. Adagietto (...)
575	6	16	The Lark
576	7	16	Adagietto from Symphony No. 5 in C-sharp minor
577	8	16	Pictures At An Exhibition: The Great Gate Of Kiev
578	9	16	Pictures At An Exhibition: Promenade I
579	10	16	Pictures At An Exhibition - Orchestrated By Ma...

This example is quite informative about the strengths and limitations of

Generating Recommendations that Promoted a Selected Group of Artists

The goal of the modified recommendations was to promote artists in a selected group; here, that group was the approximately 10,000 female mid-tier artists enumerated in the table PromotedArtist. The baseline recommendations are generated by starting from the visible tracks on a survey playlist, then joining these tracks on the full TrackPair table, and predicting the top 10 tracks that are not duplicates. Each one of the resulting track frequencies can thus be viewed as an un-normalized probability in the Naive Bayes model. The idea behind boosting the promoted artists was very simple. We set a policy amplification rate of 2.0. Every artist that was part of the promoted group had their frequency muliplied by this factor. These adjusted frequencies were then used to generate the recommendations. Here is the SQL code that does this:

11_v_SurveyRecommendPromoted.sql / Recommend Tracks Favoring Promoted Artists Based on One Survey Playlist

```

DROP VIEW IF EXISTS v.SurveyRecommendationsPromoted;
GO

CREATE VIEW v.SurveyRecommendationsPromoted
AS

WITH t1 AS (
SELECT
    -- Integer IDs
    sp.SurveyPlaylistID,
    tr.TrackID,
    -- Description of the Playlist
    sr.RecipientName,
    sr.RecipientEmail,
    sp.PlaylistName,
    sp.PlaylistNum,
    -- Recommendations
    row_number() OVER
        (partition BY sp.SurveyPlaylistID
         ORDER BY pr.Frequency * COALESCE(par.PromotionFactor, 1.0) DESC)
        AS Position,
    pr.Frequency * COALESCE(par.PromotionFactor, 1.0) AS AdjustedFrequency,
    -- The recommended track
    tr.TrackName,
    tr.TrackUri
FROM
    -- Start with all survey predictions
    dbo.SurveyPrediction AS pr
    -- The survey playlist
    INNER JOIN dbo.SurveyPlaylist AS sp ON
        sp.SurveyPlaylistID = pr.SurveyPlaylistID
    -- The track predicted
    INNER JOIN dbo.Track AS tr ON
        tr.TrackID = pr.TrackID
    -- The PromotedArtist on this track if applicable
    LEFT JOIN dbo.PromotedArtist AS par ON
        par.ArtistID = tr.ArtistID
    -- The survey response
    INNER JOIN dbo.SurveyResponse AS sr ON
        sr.SurveyResponseID = sp.SurveyResponseID
)
SELECT
    t1.SurveyPlaylistID,
    t1.TrackID,
    t1.RecipientName,
    t1.RecipientEmail,
    t1.PlaylistName,
    t1.PlaylistNum,
    t1.Position,
    t1.AdjustedFrequency,
    t1.TrackName,
    t1.TrackUri,
    row_number() OVER (ORDER BY t1.SurveyPlaylistID, t1.Position)
    AS SortOrder
FROM
    t1
WHERE
    t1.Position <= 10;

```

The database view above is using a different view called v.SurveyPredictions that is essentially replicating the predictions in the

TrackPair model using the survey playlists as input. The only interesting wrinkle here is the left join onto the PromotedArtist table. This is where the promotion factor of 2.0 is picked up. The `coalesce(par.PromotionFactor, 1.0)` has the effect of leaving non-promoted artists unchanged. This design could allow different promotion rates for each artist. In practice they were all set to 2.0.

Effect of Promoting Selected Artists on their Recommendation Rate

We can estimate the prevalence of recommending promoted artists before and after this change. This was estimated across the entire data set rather than the surveys, because the survey data set was small. The overall rate with which female midtier artists appeared in the top 100 recommended tracks was approximately **0.8576%** before any adjustment was applied. Applying a 2.0x factor to the frequency of these tracks increased their appearance rate to **1.3209%** in the top 100 tracks, i.e. increased them by a factor of 1.54. This factor was not the same as the 2.0x factor applied to the frequencies, because increasing a frequency is not a linear mapping to appearing in the top 100. To take an extreme example, if an obscure track had a frequency of 1, it would never appear in the top 100 even if we multiplied its frequency by huge promotion factor like 100.

The small absolute size of the female midtier artists is an indication that measuring an effect from this intervention would be very challenging even if we had a larger sample size. When we designed the study, we specified the largest possible set of promoted artists to mitigate this effect. In particular, artists on Spotify are ranked into 5 tiers 1 through 5, with 1 being the most popular and 5 being the most obscure. The initial proposal was to promote female artists in tiers 3-5, but we modified that to include tiers 2-5 to make the promoted group larger. Even with that change, they are representing a very small slice of the tracks people put on playlists. The underlying reality is that the music industry is very much a "winner take all" type of business, with the top artists (in tier 1) getting the lion's share of both the money and listening time.